PROJECT SUMMARY

Testing systematic predictions about future events against observed outcomes is generally seen as the most formidable validity check of statistical and theoretical models. Yet, political scientists rarely make predictions about the future. Empirical models are seldom applied to out-of-sample data much less used to make predictions about future outcomes. Instead, researchers focus on developing and validating theories that explain past events.

In part, this results from the fact that it is difficult to make accurate predictions about complex social phenomena. Yet, research in political science could gain immensely in its policy relevance if predictions were more common and more accurate. Improved forecasting of important political events would make research more germane to policymakers and the general public who are less interested in explaining the past than anticipating the future. From a scientific standpoint, greater attention to forecasting would facilitate stringent validation of theoretical and statistical models since truly causal models perform better in out-of-sample forecasting.

We propose to extend a promising statistical method – ensemble Bayesian model averaging (EBMA) – and to develop software that will aid researchers across disciplines to make more accurate forecasts. This project builds on work in the fields of meteorology, fluid dynamics, and statistics to (1) extend the method for application to a wider array of outcomes (e.g., binary data), (2) provide freely available software that implements both maximum likelihood and Bayesian estimation techniques, and (3) publish papers that provide accessible explanations of the method and real-world social science applications.

In essence, EBMA makes more accurate predictions possible by pooling predictions from multiple forecast models. The weight assigned to each component is calibrated via its predictive performance in some training period. These component models can be diverse. They need not share covariates, functional forms, or error structures. Indeed, the components may not even be statistical models, but may be predictions generated by agent based models, stochastic simulations, or subject-matter experts. EBMA pools information from heterogenous sources to improve forecast accuracy of outcomes determined by complex processes. In practice, the method provides superior predictive power relative to any single component. Further, each component model's weight can be interpreted as the posterior probability that this particular model reflects the "true" data generating process and thus facilitates meaningful model comparison.

Intellectual merit: This project will address several substantive questions in political science with a focus on the field of international relations, where policymakers have a particular demand for improved forecasting. We will also further develop statistical techniques that will improve the capabilities of researchers across disciplines to produce more accurate forecasts of future events. The methodological advancement made in this project will be available to the larger scholarly community and influence research agendas in multiple fields. EBMA has received considerable attention in the fields of statistics, meteorology, and (to a lesser extent) economics. Yet, it has not been advanced in the methodological directions we are proposing here. Additionally, our work will make EBMA available to a wider audience of researchers across the physical and social sciences. This is particularly important as existing research projects and software packages are narrowly tailored to the needs of weather forecasting.

Broader impacts: The principal investigators are active members of the research community at the intersection of statistics and the social sciences. Their work is widely read in political science and other disciplines. Further, at least two graduate student in political science will be included in the research and will gain experience in large-scale research projects.

PROJECT DESCRIPTION

Testing systematic predictions about future events against observed outcomes is generally seen as the most formidable validity check of statistical and theoretical models. Yet, political scientists rarely make predictions about the future. Empirical models are seldom applied to out-of-sample data much less used to make predictions about future outcomes. Instead, researchers focus on developing and validating theories that explain past events.

In part, this results from the fact that it is difficult to make accurate predictions about complex social phenomena. Yet, research in political science could gain immensely in its policy relevance if predictions were more common and more accurate. Improved forecasting of important political events would make research more germane to policymakers and the general public who are less interested in explaining the past than anticipating and altering the future. From a scientific standpoint, greater attention to forecasting would facilitate stringent validation of theoretical and statistical models since truly causal models perform better in out-of-sample forecasting.

We propose to extend a promising statistical method – ensemble Bayesian model averaging (EBMA) – and to develop software that will aid researchers across disciplines to make more accurate forecasts. This project builds on work in the fields of meteorology, fluid dynamics, and statistics to (1) extend the method for application to a wider array of outcomes (e.g., binary data), (2) provide freely available software that implements both maximum likelihood and Bayesian estimation techniques, and (3) publish papers that provide accessible explanations of the method and real-world social science applications.

First, we briefly review existing political science research aimed at forecasting. We then present the EBMA method. Next we present results from empirical applications of the method to the areas of insurgency events on the Pacific Rim, U.S. presidential elections, and voting on the U.S. Supreme Court. We discuss our proposed research agenda and conclude with a discussion of results from prior NSF supported research.

1. DYNAMIC FORECASTING IN POLITICAL SCIENCE

Although forecasting is a rare exercise in political science, there are an increasing number of exceptions. In most cases, "forecasts" are conceptualized as an exercise in which the predicted values of a dependent variable are calculated based on a specific statistical model and then compared with observed values (e.g., Hildebrand, Laing and Rosenthal 1976). In many instances, this becomes nothing more than an analysis of residuals. In others, the focus is on randomly selecting subsets of the data to be excluded during model development for cross-validation. However, there is also a more limited tradition of making predictions about events that have not yet occurred.

An early proponent of using statistical models to make predictions in the realm of international relations (IR) was Stephen Andriole (Andriole and Young 1977). In 1978, a volume edited by Nazli Choucri and Thomas Robinson provided an overview of the then current work in forecasting in IR. Much of this work was done in the context of policy oriented research for the U.S. government during the Vietnam War. Subsequently, there were a variety of efforts to create or evaluate forecasts of international conflict including Freeman and Job (1979), Singer and Wallace (1979), and Vincent (1980). In addition, a few efforts began to generate forecasts of domestic conflict (e.g., Gurr and Lichbach 1986). Recent years, however, have witnessed increasing interest in prediction across a wide array of contexts in IR.[1] The 2011 special issue of *Conflict Management and Peace Science*

---

[1]An incomplete list of recent work would include Krause (1997), Davies and Gurr (1998), Pevehouse and Goldstein (1999), Schrodt and Gerner (2000), King and Zeng (2001), O'Brien (2002), de Mesquita (2002), Fearon and Laitin (2003), De Marchi, Gelpi and Grynaviski (2004), Ruger et al. (2004), Enders and Sandler (2005), Leblang and

on predicting conflict in the field of IR exemplifies the growing emphasis on forecasting within that subfield (c.f., Schneider, Gleditsch and Carey 2011; Mesquita 2011; Brandt, Freeman and Schrodt 2011).

Outside of IR, forecasting in political science has largely taken place in the context of election research. In the 1990s, scholars of U.S. politics began publishing predictions of presidential elections (Campbell and Wink 1990; Campbell 1992). These efforts were anticipated by the efforts of several economists, most notably the forecasts established by Ray C. Fair (1978). As we discuss below, predicting U.S. presidential and congressional elections has since developed into a regular exercise. Moreover, similar efforts have begun to forecast election outcomes in France (e.g., Jerome, Jerome and Lewis-Beck 1999) and the United Kingdom (e.g., Whiteley 2005).[2]

While efforts to predict future outcomes remains uncommon, research that combines multiple forecasts are nearly non-existant. To our knowledge, the only example is the PollyVote project (c.f. Graefe et al. 2010), which combines multiple predictions to forecast U.S. presidential elections.

EBMA is a statistical approach that promises to improve prediction of complex social outcomes and build upon the increased interest in generating political forecasts. In essence, EBMA makes more accurate predictions possible by pooling information from multiple forecast models. The weight assigned to each component is calibrated via its predictive performance in some training period. These component models can be diverse. They need not share covariates, functional forms, or error structures. Indeed, the components may not even be statistical models, but may be predictions generated by agent based models, stochastic simulations, or subject-matter experts. EBMA pools information from heterogenous sources to improve forecast accuracy of outcomes determined by complex processes.

Such ensemble methods have been shown to significantly reduce prediction error in two important ways. First, across subject domains, ensemble predictions are usually more accurate than any individual component model. Second, they are significantly less likely to make dramatically incorrect predictions (Armstrong 2001; Raftery et al. 2005). Combining forecasts not only reduces reliance on single data sources and methodologies (which lowers the likelihood of dramatic errors), but also allows for the incorporation of more information than any one theoretical or statistical model is likely to include in isolation. In the next section, we present the details of the EBMA method that we propose to extend for social science research with support from the NSF.

2. Ensemble Bayesian model averaging

As the previous section has shown, while predictive models are still underutilized, an increasing number of scholars have developed forecasting models for specific research domains. As the number of forecasting models proliferate, however, there is a growing need to develop methods to pool across models and methodologies to generate more accurate forecasts. Very often, specific predictive models prove to be correct only for certain subsets of observations. Moreover, as we show in our examples below, specific models tend to be more fragile and sensitive to unusual events or particular data issues than ensemble methods.

The research proposed here will aid the newfound emphasis on prediction by advancing recent statistical research aimed at integrating multiple predictions into a single improved forecast. In

Satyanath (2006), Ward, Siverson and Cao (2007), Brandt, Colaresi and Freeman (2008), Bennett and Stam (2009), and Gleditsch and Ward (2010). A summary of classified efforts is reported in Feder (2002). An overview of some of the historical efforts along with a description of current thinking about forecasting and decision-support is given by O'Brien (2010).

[2]Lewis-Beck (2005) provides a more in-depth discussion of election forecasting in a comparative context.

particular, we propose adapting an ensemble method first developed for application to the most developed prediction models in existence – weather forecasting models. To generate predictive distributions of outcomes (e.g., temperature), weather researchers apply ensemble models to forecasts generated from multiple models (Raftery et al. 2005). State-of-the-art ensemble forecasts aggregate multiple runs of (often multiple) weather prediction models into a single forecast.

The EBMA method we plan to extend, first proposed by Raftery et al. (2005), pools across various forecasts, while meaningfully incorporating *a priori* uncertainty about the "best" model. It assumes that no particular model or forecasting method can fully encapsulate the true data generating process. Rather, various research teams or statistical techniques will reflect different facets of reality. EBMA collects *all* of the insights from multiple forecasting efforts in a coherent manner. The aim is not to choose some "best" model, but rather to incorporate the insights and knowledge implicit in various forecasting efforts via statistical post-processing.

EBMA itself is an extension of the Bayesian Model Averaging (BMA) methodology (c.f., Madigan and Raftery 1994; Draper 1995; Raftery 1995; Hoeting et al. 1999; Clyde 2003; Raftery and Zheng 2003; Clyde and George 2004), which has received considerable attention in the field of statistics. BMA was first introduced to political science by Bartels (1997), and has been applied in a number of contexts (e.g., Bartels and Zaller 2001; Gill 2004; Imai and King 2004; Geer and Lau 2006). Montgomery and Nyhan (2010) provide a more in-depth discussion of BMA and its applications in political science.

2.1. **Mathematical intuition.** The basic BMA approach to forecasting is as follows. Assume we have some quantity of interest in the future to forecast, $\mathbf{y}^*$, based on previously collected training data $\mathbf{y}^T$ that is fit to K statistical models, $M_1, M_2, \ldots, M_K$. Each model, M_k, is assumed to come from the prior probability distribution $M_k \sim \pi(M_k)$, and the probability distribution function (PDF) for the training data is $p(\mathbf{y}^T|M_k)$. The outcome of interest is distributed $p(\mathbf{y}^*|M_k)$. Applying Bayes Rule, we get that

$$(1) \qquad p(M_k|\mathbf{y}^T) = \frac{p(\mathbf{y}^T|M_k)\pi(M_k)}{\sum\limits_{k=1}^{K} p(\mathbf{y}^T|M_k)\pi(M_k)}.$$

and the marginal predictive PDF for y^* is

$$(2) \qquad p(\mathbf{y}^*) = \sum_{k=1}^{K} p(\mathbf{y}^*|M_k)p(M_k|\mathbf{y}^T).$$

The BMA PDF (2) can be viewed as the weighted average of the component PDFs where the weights are determined by each model's performance within the training data. Likewise, we can make a deterministic estimate as the weighted predictions of the components, denoted

$$(3) \qquad E(\mathbf{y}^*) = \sum_{k=1}^{K} E(\mathbf{y}^*|M_k)p(M_k|\mathbf{y}^T).$$

2.2. **EBMA for dynamic settings.** We now turn to applying this basic BMA technology to prediction in a dynamic setting. In generating predictions of important events (e.g., domestic crises or international disputes), the task is to first build a statistical model for some set of countries S in time periods T, which we refer to as the training period.[3] Using the same statistical model (or

---

[3]Sloughter et al. (2007) make predictions for only one future time period, and use only a subset of past time-periods (they recommend 30) in their training data. Thus, predictions are made sequentially with the entire EBMA procedure being re-calculated for each future event as observations are moved from the out-of-sample period T^* into the training

general technique in the case of expert predictions), we then generate forecasts, $\mathbf{f}_k$, for observations S in future time periods T^*.

Let us assume that we have K forecasting models predicting outcome events $\mathbf{y}$. Each component forecast, $\mathbf{f}_k$, is associated with a component PDF, $g_k(\mathbf{y}|\mathbf{f}_k)$, which may be the original predictive PDF from the forecast model or a bias corrected forecast. These components are the conditional PDFs of outcome $\mathbf{y}$ given the kth forecast, $\mathbf{f}_k$, conditional on it being the "best" forecast in the ensemble. For example, the posterior PDF of an outcome y_{st} for some country $s \in S$ in period $t \in T$ given the forecast f_{kst} from model k is $g_k(y_{st}|f_{kst})$. This assumes that $P(M_k|\mathbf{y}) \equiv w_k = 1$, or that the posterior odds of model k is unity.

The EBMA PDF is then a finite mixture of the K component forecasts, denoted

$$(4) \qquad p(\mathbf{y}|\mathbf{f}_1, \ldots, \mathbf{f}_K) = \sum_{k=1}^{K} w_k g_k(\mathbf{y}|\mathbf{f}_k),$$

where the weight w_k is based on forecast k's relative predictive performance in the training period T. The w_k's $\in [0,1]$ are probabilities and $\sum_{k=1}^{K} w_k = 1$. The specific PDF of for an event y_{st^*} will then be

$$(5) \qquad p(y_{st^*}|f_{1st^*}, \ldots, f_{Kst^*}) = \sum_{k=1}^{K} w_k g_k(y_{st^*}|f_{kst^*}).$$

2.3. **EBMA for normally distributed outcomes.** When forecasting outcomes that are distributed according to the normal distribution, Raftery et al. (2005) propose approximating the conditional PDF as a normal distribution centered at a linear transformation of the individual forecast, $g_k(\mathbf{y}|\mathbf{f}_k) = N(a_{k0} + a_{k1}\mathbf{f}_k, \sigma^2)$. Using (4) above, the EBMA PDF is then

$$(6) \qquad p(\mathbf{y}|\mathbf{f}_1, \ldots, \mathbf{f}_K) = \sum_{k=1}^{K} w_k N(a_{k0} + a_{k1}\mathbf{f}_k, \sigma^2).$$

2.4. **The dichotomous outcome model.** Past work on EBMA does not apply directly to the prediction of many political events because the assumed PDFs are normal, Poisson, or gamma. In many settings (e.g., international conflicts), the data are not sufficiently fine-grained to justify these distributional assumptions. Usually, the outcomes of interest are dichotomous indicators for whether an event (e.g., civil war) has occurred in a given time period in a specified country or region. Thus, none of the distributional assumptions used in past work are appropriate in this context. Fortunately, it is possible to extend Sloughter et al. (2007) and Sloughter, Gneiting and Raftery (2010) to deal appropriately with binary outcomes.

We follow Sloughter et al. (2007) and Hamill, Whitaker and Wei (2004) in using logistic regression after a power transformation of the forecast to reduce prediction bias. For notational ease, we assume that $\mathbf{f}_k$ is the forecast after the adjustment for bias reduction. Therefore, let $\mathbf{f}'_k \in [0,1]$ be the forecast on the predicted probability scale and

$$(7) \qquad \mathbf{f}_k = \left[(1 + \mathrm{logit}(\mathbf{f}'_k))^{1/b} - 1 \right] I \left[\mathbf{f}'_k > \frac{1}{2} \right] - \left[(1 + \mathrm{logit}(|\mathbf{f}'_k|))^{1/b} - 1 \right] I \left[\mathbf{f}'_k < \frac{1}{2} \right],$$

---

set T. Another alternative is to simply divide *all* the data into discrete training and test periods for the entire procedure. We use both approaches in our examples below.

where $I[.]$ is the general indicator function. Hamill, Whitaker and Wei (2004) recommend setting $b = 4$, while Sloughter et al. (2007) use $b = 3$. We found that $b = 4$ works best in the examples below, but other analysts may try alternative specifications.[4]

The logistic model for the outcome variables is

$$(8) \qquad \text{logit } P(\mathbf{y} = 1|\mathbf{f}_k) \equiv \log \frac{P(\mathbf{y} = 1|\mathbf{f}_k)}{P(\mathbf{y} = 0|\mathbf{f}_k)} = a_{k0} + a_{k1}\mathbf{f}_k.$$

The conditional PDF of some within-sample event, given the forecast f_{kst} and the assumption that k is the true model, can be written:

$$(9) \qquad g_k(y_{st}|f_{kst}) = P(y_{st} = 1|f_{kst})I[y_{st} = 1] + P(y_{st} = 0|f_{kst})I[y_{st} = 0].$$

Applying this to (4), the PDF of the final EBMA model for y_{st} is

$$(10) \qquad p(y_{st}|f_{1st}, f_{2st}, \ldots, f_{Kst}) = \sum_{k=1}^{K} w_k[P(y_{st} = 1|f_{kst})I[y_{st} = 1] + P(y_{st} = 0|f_{kst})I[y_{st} = 0]].$$

2.5. Parameter estimation by maximum likelihood and EM algorithm.

Parameter estimation is conducted using only the data from the training period T. The parameters a_{0k} and a_{1k} are specific to each individual component model. For model k, these parameters can be estimated as traditional linear models where $\mathbf{y}$ is the dependent variable and the covariate list includes only $\mathbf{f}_k$.

The difficulty is in estimating the weighting parameters, $w_k \ \forall \ k \in [1, 2, \ldots, K]$. One approach we propose to implement with NSF support is to place priors on all parameters and conduct a fully Bayesian analysis with Markov chain Monte Carlo techniques (c.f. Vrugt, Diks and Clark 2008). For the moment, however, we have followed Raftery et al. (2005) and Sloughter et al. (2007) in using maximum likelihood methods.

With standard independence assumptions, the log-likelihood for the model weights is

$$(11) \qquad \ell(w_1, \ldots, w_K|a_{01}, \ldots, a_{0K}; a_{11}, \ldots, a_{1K}) = \sum_{s,t} \log p(y_{st}|f_{1st}, \ldots, f_{Kst}).$$

where the summation is over values of s and t that index all observations in the training time period, and $p(y_{st}|f_{1st}, \ldots, f_{Kst})$ is given by (10). The log-likelihood function cannot be maximized analytically, but Raftery et al. (2005) suggests using the expectation-maximization (EM) algorithm. We introduce the unobserved quantities z_{kst}, which represent the posterior probability of model k for observation y_{st}. The E step involves calculating estimates for these unobserved quantities using the formula

$$(12) \qquad \hat{z}_{kst}^{(j+1)} = \frac{\hat{w}_k^{(j)} p^{(j)}(y_{st}|f_{kst})}{\sum_{k=1}^{K} \hat{w}_k^{(j)} p^{(j)}(y_{st}|f_{kst})},$$

where the superscript j refers to the jth iteration of the EM algorithm.

It follows that $w_k^{(j)}$ is the estimate of w_k in the jth iteration and $p^{(j)}(.)$ is shown in (10). Assuming these estimates of z_{kst} are correct, it is then straight forward to derive the maximizing value for the model weights. Thus, the M step estimates these as $\hat{w}_k^{(j+1)} = \frac{1}{n}\sum_{s,t} \hat{z}_{kst}^{(j+1)}$, where n represents the number of observations in the training dataset.[5]

---

[4] This transformation dampens the effect of extreme observations and reduce over-fitting.

[5] In the case of normally distributed data, $\hat{\sigma}^{2(j+1)} = \frac{1}{n}\sum_{s,t}\sum_{k=1}^{K} \hat{z}_{kst}^{(j+1)}(y_{st} - f_{kst})^2$.

The E and M steps are iterated until the improvement in the log-likelihood is no larger than some pre-defined tolerance. Although the log-likelihood will increase after each iteration of the algorithm, convergence is only guaranteed to a local maximum of the likelihood function. Convergence to the global maximum is not assured, and the model is therefore sensitive to initial conditions.[6] In the examples below, we begin with the assumption that all models are equally important, $w_k = \frac{1}{K} \, \forall \, k \in [1, \ldots, K]$.

2.6. **Ensemble prediction.** With these parameter estimates completed, it is now possible to generate ensemble forecasts. If our forecasts, $\mathbf{f}_k$, are generated from a statistical model, we now generate a new set of predictions f_{kst^*} from the previously fitted models. For convenience, let $\hat{\mathbf{a}}_k \equiv (\hat{a}_{k0}, \hat{a}_{k1})$. For some dichotomous observation in country $s \in S$ in some out-of-sample period $t^* \in T^*$, we can see that

$$(13) \quad P(y_{st^*} = 1 | f_{1st^*}, \ldots, f_{Kst^*}; \hat{\mathbf{a}}_1, \ldots, \hat{\mathbf{a}}_K; \hat{w}_1, \ldots, \hat{w}_K) = \sum_{k=1}^{K} \hat{w}_k \text{logit}^{-1} \left(\hat{a}_{k0} + \hat{a}_{k1} f_{kst^*} \right).$$

3. EMPIRICAL APPLICATIONS

3.1. **Application to insurgency forecasting.** Our first example applies the EBMA methodology to data collected for the Integrated Crisis Early Warning Systems (ICEWS) project, sponsored by the Defense Advanced Research Projects Agency (DARPA). The task of the ICEWS project is train models on data (focusing on five outcomes of interest) for 29 countries for every month from 1997 through the present and to then make accurate predictions about expected crisis events over the subsequent three months.[7] For purposes of demonstration we focus on only one of these outcomes – violent insurgency.

The bulk of this data is gleaned from natural language processing of a continuously updated harvest of news stories, primarily taken from Lexus/Nexus and Factiva archives. These are processed with a version of the TABARI processor for events developed by Philip Schrodt and colleagues in the context of the event data project and are described in more detail elsewhere (see `http://eventdata.psu.edu/`). These data are augmented with a variety of covariates. In particular, we use attributes (coded on a monthly or yearly basis) from the Polity and World Bank datasets. We also include information about election cycles (if any), events in neighboring countries, and the length of shared borders.

3.1.1. *Component Models.* In the remainder of this sub-section we apply the EBMA method to make predictions for the occurrence of insurgency in these 29 countries. We fit three exemplar models using data for the in-sample period ranging from January 1999 to December 2008 and fit an EBMA model. We then make out-of-sample predictions for the period from January 2009 to December 2010 for the component and EBMA forecasts. To provide variation in the complexity (as well as accuracy) of the component models, we included the following models:

---

[6]In future research, we propose to explore these convergence issues more fully with special attention paid to comparison with fully Bayesian implementations.

[7]The twenty-nine countries are Australia, Bangladesh, Bhutan, Cambodia, China, Comoros, Fiji, India, Indonesia, Japan, Laos, Madagascar, Malaysia, Mauritius, Mongolia, Myanmar, Nepal, New Zealand, North Korea, Papua New Guinea, Philippines, Russia, Singapore, Solomon Islands, South Korea, Sri Lanka, Taiwan, Thailand, and Vietnam. This set is not a random sample, but rather constitutes the countries of population greater than $500,000$ that are in the Area of Responsibility of the US Pacific Command.

- **SAE**: This is one model developed as part of the ICEWS project and was designed by Strategic Analysis Enterprises. The model is a generalized linear model using 27 different independent variables[8]. All of variables are taken from the ICEWS event-stream data.
- **GLM**: For the purposes of demonstrating the properties of the EBMA method, we built a crude logistic model that includes only *population size* and *GDP growth* (both lagged three months).
- **LMER**: This is a generalized linear mixed effects model including random country-level intercepts. In addition to the variables from the GLM model, the list of covariates includes: the *executive constrain* and *competitiveness of participation* variables from the Polity IV dataset (Marshall, Jaggers and Gurr 2009), *population size*, *proximity to election*,[9] and a *spatial lag* that reflects recent occurrences of insurgencies in the countries' geographic neighbors.[10]

3.1.2. *Results.* Table 1 shows the EBMA model parameters and fit statistics associated with the component models and the EBMA forecasts generated for the in-sample time period. The first column shows the weights that the EBMA model assigned each component. As can be seen, the GLM model is effectively excluded, while the SAE model carries the greatest weight followed by the LMER model. The constant term associated with each component corresponds to the term a_{k0} in Equation 8, while predictor term corresponds to a_{k1}.

TABLE 1. In-sample results: Estimated model weights, parameters, and fit statistics for EBMA deterministic forecast and all component forecasts of insurgency in 29 countries of the Pacific rim.

	Weight	Constant	Predictor	AUC	PRE	Brier	% Correct
SAE	0.57	0.04	7.46	0.96	0.48	0.04	94.11
LMER	0.43	6.08	28.25	0.96	0.01	0.07	88.79
GLM	0.00	0.57	8.16	0.65	0.00	0.10	88.65
EBMA				0.97	0.55	0.04	94.94
n=3,480							

The other columns in Table 1 are fit statistics. AUC is the area under the Receiver-Operating Characteristic (ROC) curve. The advantage of using ROC curves is that it evaluates forecasts in a way that is less dependent on an arbitrary cutoff point. A value of 1 would mean that all observations were predicted correctly at all possible cutoff points (King and Zeng 2001).

We compare the models using three additional metrics. The proportional reduction in error (PRE) is the percentage increase of correctly predicted observations relative to some pre-defined base model. The base model in this case is predicting "no insurgencies" for all observations. Insurgencies are relatively rare events. Thus, predicting a zero for all observations leads to an 89% correct prediction rate. The Brier score is the average squared deviation of the predicted probability from the true event (0 or 1). Thus, a lower score corresponds to higher forecast accuracy (Brier 1950). In addition, we calculate the percentage of observations that each model would predict correctly using a 0.5 threshold on the predicted probability scale.

There are two aspects of Table 1 that are important to note. First, the EBMA model does at least as well (and usually better) than all of the component models on each our model fit statistics. The EBMA model has the highest AUC, PRE, and % correct. In addition, it is tied for the lowest Brier score with the SAE model. Second, in this example the EBMA procedure assigns probability

---

[8]See `strategicanalysisenterprises.com` for more details

[9]This is calculated as the number of days to the next or from the last election, whichever is closer.

[10]Geographical proximity is measured in terms of the length of the shared border between the two countries.

TABLE 2. Out-of-sample results: Fit statistics for EBMA deterministic forecast and all component model forecasts of insurgency in 29 countries of the Pacific rim.

	AUC	PRE	Brier	% Correct
SAE	0.96	0.04	0.06	89.80
LMER	0.97	0.00	0.07	89.37
GLM	0.84	0.00	0.09	89.37
EBMA	0.96	0.18	0.05	91.24
n=696				

weights to components according to their in-sample performance. The highest model weight (0.57) is assigned to the SAE model, which appears to be the best (or tied for the best) on all of our fit statistics. Meanwhile, the lowest weight (0.00) is assigned to the rudimentary GLM model.

FIGURE 1. Separation plots for in-sample predictions of the ICEWS data (n=3,480). For each model, observations are shown from left to right in order of increasing predicted probability of insurgency (shown as the black line). Observations where insurgency actually occurred are shown in red.


Figure 1 shows separation plots for the EBMA model and the individual components (Greenhill, Ward and Sacks 2011). In each plot, the observations are ordered from left to right by increasing predicted probabilities of insurgency (as predicted by the particular model). The black line corresponds to the predicted probability produced by the relevant model for each observation and actual occurrences of insurgencies are colored red. Figure 1 shows visually that the GLM model perform very poorly, whereas of the SAE model is the best component. More importantly, the overall best performance is the EBMA forecast. The separation plots show that the EBMA model produces few false positives and even fewer false negatives than any of the component models.

The more interesting evaluation of the EBMA method is its out-of-sample predictive power. Table 2 shows fit statistics for the individual components as well as the EBMA forecasts for observations in the 24 months following the training period. While the EBMA model has a marginally smaller area under the ROC curve than the LMER models, it outperforms all component models on the other fit statistics. In particular, the EBMA model has the highest PRE at 0.18. Since it is possible to predict 89.22% of these observations correctly by forecasting no insurgency, an 18% reduction of error relative to the baseline model is quite substantial.

Figure 2 shows the separation plots for the components as well as the EBMA forecasts for the out-of-sample data. The EBMA model performs better than any of the individual components, with high predicted probabilities for most observed insurgencies. Taking both the fit statistics

and the visual evidence together, we can conclude that the EBMA model leads to a substantial improvement in out-of-sample forecasts relative to its components.

FIGURE 2. Separation plots for out-of-sample predictions of the ICEWS data (n=696). For each model, observations are shown from left to right in order of increasing predicted probability (shown as the black line). Observations where insurgency actually occurred are shown in red.

3.2. **Application to US presidential election forecasts.** For the past several presidential elections, a number of research teams have developed forecasting models and published their predictions in advance of the election. For example, before the most recent election, a symposium of forecasts was published in *PS: Political Science and Politics* with forecasts of presidential and congressional vote shares developed by Campbell (2008), Norpoth (2008), Lewis-Beck and Tien (2008), Abramowitz (2008), Erikson and Wlezien (2008), Holbrook (2008), Lockerbie (2008) and Cuzàn and Bundrick (2008). Responses to the forecast and evaluations were published in a subsequent issue of the journal.[11] Predicting presidential elections has also drawn the attention of economists seeking to understand the relationship between economic fundamentals and political outcomes. Two prominent examples include work by Ray Fair (2010) and Douglas Hibbs (2000).

3.2.1. *Component Models.* In the rest of this sub-section, we replicate several of these models and demonstrate the usefulness of the EBMA methodology for improving the prediction of single important events. We include six of the most widely cited presidential forecasting models:

- **Campbell**: Joseph Campbell's "Trial-Heat and Economy Model" (Campbell 2008).
- **Lewis-Beck/Tien**: Lewis-Beck and Tien's "Jobs Model Forecast" (Lewis-Beck and Tien 2008),
- **Erikson/Wlezien**: Erikson and Wlezien's "Leading Economic Indicators and Poll" forecast,[12]
- **Fair**: Fair's presidential vote-share model,[13]
- **Hibbs**: Hibbs' "Bread and Peace Model" (Hibbs 2000),
- **Abramowitz**: and, the "Time-for-Change Model" created by Abramowitz (2008).

With the exception of the Hibbs forecast, the models are simple linear regressions. The dependent variable is the share of the two-party vote received by the incumbent-party candidate.[14]

---

[11]In 1999, an entire issue of the *International Journal of Forecasting* was dedicated to the task of predicting presidential elections (Brown and Chappell 1999).

[12]We replicated Column 2 in Table 2 from Erikson and Wlezien (2008).

[13]The model here replicates Equation 1 in Fair (2010).

[14]The data to replicate the models by Abramowitz (2008), Campbell (2008), Erikson and Wlezien (2008), and Lewis-Beck and Tien (2008) were provided to us in personal correspondence with the respective authors. The remaining data was downloaded from the personal websites of Ray C. Fair 2011 and Douglas Hibbs 2011.

3.2.2. *Results.* Rather than selecting a single training period (as in the insurgency analysis) we generate sequential predictions. For each year from 1976 to 2008, we use all available prior data to fit the component models.[15] We then fit the EBMA model using the components' in-sample performances for election years beginning with 1952 (the year when all models begin generating prediction). For example, to generate prediction for the 1988 election we used the in-sample performance of each component for the 1952-1984 period to estimate model weights.[16]

Table 3 provides results for the 2008 and 2004 elections. The first column shows the out-of-sample prediction errors, calculated as $y_{predicted} - y_{observed}$, for each component model and the EBMA forecast. Table 3 also shows the weights assigned to each model as well as the in-sample root mean squared error (RMSE) and mean absolute error (MAE) for the components and the EBMA forecasts.

TABLE 3. Prediction errors, model weights, and in-sample fit statistics for component and EBMA forecasts of the 2008 and 2004 elections. Models are sequentially fit using all prior elections.

| | 2008 Election | | | | 2004 Election | | | |
	Pred. Error	Weights	RMSE	MAE	Pred. Error	Weights	RMSE	MAE
Campbell	5.92	0.34	1.69	1.36	0.11	0.29	1.75	1.45
Lewis-Beck/Tien	-1.95	0.05	1.64	1.40	0.30	0.00	1.70	1.49
Erikson/Wlezien	-0.11	0.20	2.81	2.18	3.70	0.21	2.72	2.06
Fair	-1.49	0.00	2.23	1.83	3.41	0.10	2.12	1.71
Hibbs	-0.93	0.34	1.93	1.33	1.50	0.40	1.95	1.32
Abramowitz	-1.88	0.08	1.57	1.30	1.92	0.00	1.55	1.25
EBMA	-0.29		1.43	1.08	1.59		1.49	1.09

Examining Table 3, it is clear that there is not as clean a relationship between in-sample model performance and model weights as in the insurgency example. For instance, the weight for the Erikson/Wlezien model in 2008 is 0.20 even though it has almost the highest RMSE and MAE of any component. The diminished relationship between in-sample performance and weight is a result of high in-sample correlations forecasts. For instance, fitted values for the Abramowitz model are correlated at 0.94 with the Campbell model and at 0.96 with the Lewis-Beck/Tien model. Thus, conditioned on knowing these forecasts, the Abramowitz component provides limited additional information. In the end, the bulk of the weight within this cluster is assigned to the Campbell model – generally the best predictor.

With the 2004 and 2008 examples in mind, we now turn to the relative out-of-sample performance of the EBMA and component forecasts across the entire 1976-2008 period. Table 4 shows the RMSE and MAE statistics as well as the percentage of observations that fall within the 67% and 95% predictive credible intervals for each. The last column shows the continuous ranked probability scores (CRPS). The CRPS is the integral of the Brier score at all possible threshold values (Hersbach 2000). Scores closer to 0 indicate superior predictive power and offer a richer evaluation of predictive distributions than RMSE or MAE.

For our purposes here, the main result in Table 4 is that the EBMA model again outperform all components. The first two columns show this to be true in terms of predicted error (RMSE and

---

[15]For example, the Fair model uses data for election results beginning in 1916 while the Abramowitz model begins with data from the 1952 election.

[16]Because of the paucity of data, we constrain the predictor, denoted a_{1k} above, to one for all models.

TABLE 4. Fit statistics, observed coverage probabilities, and continuous ranked probability scores (CRPS) for sequentially generated predictions of presidential elections from 1976-2008.

	RMSE	MAE	67% Coverage	95% Coverage	CRPS
Campbell	2.35	1.65	0.22	0.44	1.45
Lewis-Beck/Tien	1.81	1.52	0.11	0.33	1.29
Erikson/Wlezien	2.41	1.88	0.22	0.33	1.68
Fair	2.51	2.16	0.00	0.11	1.91
Hibbs	2.20	1.60	0.22	0.44	1.38
Abramowitz	1.84	1.62	0.11	0.22	1.39
EBMA	1.63	1.26	0.67	0.78	0.95

MAE). Moreover, the CRPS statistics show it to also be true for the predictive PDFs and not just for the deterministic forecasts.

The coverage statistics demonstrate much better calibration of EBMA forecasts relative to its component models. For instance, the observed outcome falls within the 95% predictive credible interval for the Campbell model only four out of eleven times while the EBMA 95% predictive interval covers the observed outcome twice as often. In general, the EBMA forecasts is less sensitive to particular data issues or events and is less likely to produce wildly incorrect predictions.

An example of the kinds of problems that may arise when relying on a single forecast is seen in a close examination of the performance of the Campbell model. From 1952 to 2004, this model was one of the strongest performers. Indeed, it made the most accurate forecast of the 2004 election. However, as a result of the particularly late timing of the Republican Convention in 2008,[17] it was the only model to forecast a victory for John McCain. By relying on a wider array of data sources and methodologies, the EBMA method reduces the likelihood of such large misses.

3.3. **Application to Supreme Court Forecasting Project.** Our final application of EBMA is a re-analysis of data from the Supreme Court Forecasting Project (Ruger et al. 2004; Martin et al. 2004).[18] This example highlights the ability of EBMA to handle forecasts generated by classification trees, subject-matter experts, and other sources.

Throughout 2002-2003, a research team consisting of Andrew Martin, Kevin Quinn, Theodore Ruger, and Pauline Kim (henceforward MQRK) generated two sets of forecasts for every pending case. First, using data about case characteristics and justices' past voting patterns, MQRK developed classification trees to generate a binary forecasts for the expected vote of each justice on each case (voting to affirm the lower court opinion is coded as a 1). Second, MQRK recruited a team of 83 legal experts to make forecasts on particular cases in their specialty area. The list included academics, appellate attorneys, former Supreme Court clerks, and law school deans. MQRK attempted to recruit three expert forecasts for each case, although this was not possible for all cases.

The statistical model makes predictions for all 67 cases included in the MQRK analysis. Thus, we include the binary model predictions as one component forecast. However, the individual legal experts made predictions on only a handful of cases. Thus, it is not possible to calibrate each expert

---

[17]One of the only two variables in the Campbell model comes from polling data measured in early September. The model also relies on 2nd quarter GDP growth, which dramatically fell as the election approached due to the financial crisis. The net result of these two factors is that the Campbell model's prediction for 2008 was one of the largest mis-predictions in the dataset.

[18]Additional details about the project, replication files, as well as a complete listing of cases and expert forecasts are available at: http://wusct.wustl.edu/index.php.

and treat them all as individual "models." Instead, we treat all of the expert opinions as part of a single forecasting effort. We coded the expert forecast to be the mean expert forecast. This implies, that the expert forecast predicts a vote to affirm if a majority of experts polled for that case predict an affirming vote. We fit an EBMA model using all cases with docket numbers dating from 2001 (n=395).[19] We then made EBMA forecasts for the remaining 296 cases with 2002 docket numbers.

3.3.1. *Results.* Table 5 shows the component weights for the two forecasts and the out-of-sample fit statistics for the MQRK classification trees, subject-matter experts, and EBMA forecasts. Once again, the results show that the EBMA procedure outperforms all components (even when there are only two). In terms of AUC, Brier scores, and correct predictions, the EBMA forecast out-performs both the statistical model and the combined subject-matter experts. In addition, the EBMA forecast offers a significantly better PRE.[20]

TABLE 5. Out-of-sample results: Fit statistics for EBMA deterministic forecast and component forecasts of Supreme Court votes on cases in the 2002-2003 session with 2002 docket numbers.

	Weight	AUC	PRE	Brier	% Correct
MQRK model	0.32	0.66	-0.02	0.29	70.56
Experts	0.68	0.62	0.15	0.23	75.23
EBMA forecast		0.70	0.21	0.18	77.10
n=214					

There is a long-standing debate in may circles of the relative strengths and weaknesses of statistical models and subject-matter experts for making predictions. Models that use quantifiable measurements and widely available (if sometimes crude) data to make predictions can make egregious errors in particular cases. Some cases may be decided by forces invisible to the statistical model but obvious to experts familiar with the case. Subject-matter experts, on the other hand, can become too focused on minutia and miss larger (if more subtle) trends in the data easily recognized by more advanced methodologies. However, the EBMA technique offers a theoretically motivated way to combine the strengths of both methods, while smoothing over their relative weaknesses, to make more accurate predictions.

4. PROPOSED RESEARCH AND TIMELINE

Thus far, we have extended prior research to make EBMA applicable to binary and continuous outcomes in political science. As the above examples demonstrate, the method already increases the accuracy of predictions. However, we propose to expand this research in four ways.

Fully Bayesian estimation: MCMC estimation of EBMA models can more efficiently handle a wider variety of predictive distributions (Vrugt, Diks and Clark 2008). Standard Bayesian methods, such as data augmentation, will allow us to build a set of statistical results and computer algorithms appropriate for an array of assumed outcome distributions. Specifically, we plan to conduct basic research to develop a class of models and MCMC samplers to handle continuos, censored-continuous, binary, and event count data.[21]

---

[19]For reasons of space, we do not present the in-sample results here. These are available upon request.

[20]The baseline model here is prediction that all votes will be to reverse the lower court. This baseline model is correct for roughly 70% of the votes in the out-of-sample period.

[21]Moreover, as the number of parameters is relatively moderate, Bayesian algorithms also promise to provide fast results while eliminating concerns about false convergence to local maxima that can occur when using EM methods.

Alternative model weights: Currently, model weights are calculated based solely on the model's deterministic forecast (i.e., their point predictions). Even for continuous data, the current EBMA model assumes that within-forecast variances (σ^2) is constant. In other words, model weights do no reflect the uncertainty associated with component predictions. Applying both Bayesian and bootstrap methods, we will develop methods to adjust the posterior model weights to incorporate the predictive PDF from each component modle. Poorly calibrated models and those making extremely uncertain predictions will be penalized.

A related issue is that EBMA makes no adjustment for model complexity. That is, model weights are based solely on the components' goodness-of-fit with no effort to adjust for their generalizability. This can lead to excessive weighting of complex and over-fit models. Since component forecasts may be agent based models, stochastic simulations, multi-level models, and the like, it is necessary to go beyond merely penalizing for the number of parameters (e.g., AIC). Complexity measures must take into account functional form and other concerns. As part of our research, we will extend several proposed methods for penalizing complexity via specifying appropriate prior structures (c.f., Pitt, Myung and Zhang 2002; Pitt and Myung 2002), implement these priors in our software, and test them against existing vague priors using both simulated and real-world data.

Software: We will develop open-source software that will be publicly available. Specifically, we will produce an R package that implements both the binary outcome version we have already developed and the additional extensions just discussed. Moreover, the package will provide a more flexible interface for users interested in ensemble forecasting outside of the weather prediction community than is currently available.

Our specific goals for the software package include: S4 compliance, computationally efficient internal functions written either in C or Java, handling of multiple outcome distributional assumptions and priors with a small set of user functions, Gaussian copula techniques for handling missing data (Hoff 2007), customizable data visualizations to facilitate EBMA and component model comparisons, exemplar datasets and vignettes based on real-world applications of the method to the social and physical sciences, and built-in convergence diagnostics for all Bayesian methods compatible with the 'coda' and 'boa' packages in R. At the completion of the project, we will submit and article to the *Journal of Statistical Software* explaining both the technical details of the package and providing tutorials explaining its available features.

Dissemination of findings: Beyond providing the software and its attendant documentation, we will disseminate the results of our research in three ways. First, we will submit articles explaining the mathematical and technical details of EBMA to journals in political methodology, economics, and applied statistics. Second, we plan to develop and publish at least two extended applications of the method to topics in political science. Our particular focus here will be: (1) improved forecasting of international crisis events; (2) election forecasting with the aim of collaborating with other scholars to generate ensemble predictions for the 2012 U.S. elections. Third, we will offer a workshop on forecasting political outcomes during the first day of the 2012 meeting of the Political Methods Conference (already scheduled to be hosted by the UNC and Duke).

Project Workplan:

01/2012 – 06/2010: In this stage, we will conduct basic research into MCMC estimation of EBMA and prior structures that penalize model complexity and ensure that posterior estimates reflect uncertainty in component forecasts.

07/2012 – 06/2013: In this stage, we will develop, test, and document the software. This will also involve gathering exemplar forecasting datasets for inclusion in package vignettes. These examples will be the basis for subsequent applied research and publication.

07/2013 – 12/2013: This stage will focus on: (1) preparing software and documentation for public dissemination, improving user interfaces, and increasing the computational efficiency of internal functions, and (2) revision and submission of results for publication.

5. RESULTS FROM PRIOR NSF SUPPORT

5.1. **NSF Grants Received by Michael D. Ward During Previous 5 Years.**

(1) 0827016 ($749,970; PI's sub $150,000) AOC: The Dynamics of Secessionist Regions: Eurasian Unrecognized Quasi-States after Kosovo's Independence 10/01/2008–09/30/2011
(2) 0631531 ($400,000) Longitudinal Network Modeling of International Relations Data 11/15/2006–10/31/2009
(3) 0433927 ($650,000; PI's sub $150,000) The Dynamics of Civil War Outcomes: Bosnia and the North Caucasus 10/01/2004–09/30/2008
(4) 0417559 ($150,000) Network Modeling of International Peace and Trade Data 10/01/2004–09/30/2006
(5) 0631531 () Longitudinal Network Modeling of International Relations Data

Only one of these grants is still open (# 1), but these funds have only recently (Spring 2011) been transferred to Duke University. The most relevant grant is # 5, which is discussed below.

Summary of Findings: One of our primary findings is that standard hazard regression methods for longitudinal relational data, using variants of proportional hazards models, are unable to properly account for temporal or relational dependence in IR data. We have explored several approaches to modeling the longitudinal dependencies. One models temporal correlations directly as a network. A second uses the temporal evolution of the latent network as a means of imparting dynamic structure into the estimation of the network parameters. A final approach, which we are completing now, involves modeling a separate time-series regression for each pair of countries, but using a special array-variate hierarchical model to allow for similarity in trade patterns across groups of countries. We show that the regularized estimates from this hierarchical model outperform existing methods in out-of-sample prediction of longitudinal trade data.

Broader Impacts: We have presented the research at a number of conferences and in-departmental seminars in the fields of statistics, biostatistics, political science, and geography. We have created the first installment of a series of open-source software packages for the analysis of relational data. These packages are widely used in the social science community conducting network analyses. Finally, we constructed a database of trade and conflict that can be accessed by our software.

Within the first year we completed the following: (1) trained a graduate student in the analysis of longitudinal relational data; (2) constructed and analyzed databases on longitudinal international relations data (including data on conflicts, trade, currency exchange, membership in IGOs); (3) developed new statistical methodologies for multivariate and longitudinal data; (4) conducted basic research in the area of multivariate statistical models, including methods related to copula modeling and reduced-rank matrix models; (5) conducted basic research into the analysis of array data, such as longitudinal trade data, using random-effects versions of multiway array methods such as PARAFAC, and developed an extension of the multivariate and matrix-variate normal distribution appropriate for modeling multiway array data.

Finally, two students learned how to gather and organize data, write technical documents, and perform independent research. Both students completed their Ph.D. in the summer of 2010. One student (John Ahlquist) received two national awards for his dissertation.

Publications Resulting from Award:

- Ward, Michael D., Randolph M. Siverson, and Xun Cao. 2007. "Disputes, Democracies, and Dependencies: A Re-examination of the Kantian Peace". *American Journal of Political Science* 51(3):583-601.
- Ward, Michael D. and Peter D. Hoff. 2008. Analyzing Dependencies in Geo-Politics and Geo-Economics. In *Contributions to Conflict Management, Peace Economics, and Development, Volume 6, War, Peace and Security*, ed. Jacques Fontanel and Manas Chatterji. Bingley, UK: Emerald Publishing, pp. 133-160.
- Krivitsky, Pavel, Mark Handcock, Adrian Raftery and Peter Hoff. 2009. "Representing Degree Distributions, Homophily and Clustering in Social Networks with Latent Cluster Models". *Social Networks* 31(3):204-213.
- Hoff, Peter. 2008. "Modeling Homophily and Stochastic Equivalence in Symmetric Relational Data". *Advances in Neural Information Processing Systems* 20: 667-674.
- Hoff, Peter. 2009. "Simulation of the Matrix Bingham-von-Mises-Fisher Distribution, With Applications to Multivariate and Relational Data". *Journal of Computational and Graphical Statistics* 18(2):438–456.
- Hoff, Peter. 2009. "A hierarchical eigenmodel for pooled covariance estimation". *Journal of the Royal Statistical Society: Series B (Statistical Methodology)* 71(5):971–992.
- Ward, Michael D. 2010. Statistical Analysis of International Interdependencies. In *The International Studies Encyclopedic Compendium: Scientific Studies of International Processes, Volume 10*, ed. Paul F. Diehl and James D. Morrow, pp. 6615-6628.
- Bakke, Kristin M., Xun Cao, John O'Loughlin and Michael D. Ward. 2009. "Social Distance in Bosnia and the North Caucasus Region of Russia: Inter- and intra-ethnic attitudes and identities". *Nations and Nationalism* 15(2):227-253.
- Hoff, Peter. 2011. "Hierarchical multilinear models for multiway data". *Computational Statistics and Data Analysis* 55(1):530–543.
- Hoff, Peter and Xiaoyue Niu. 2010. "A covariance regression model". *Statistica Sinica* Submitted.
- Ahlquist, John S. 2008. "Building and Using Strategic Capacity: Labor Union Confederations and Economic Policy". Thesis. Published. Bibliography: Doctoral Dissertation, University of Washington.
- Raftery, Adrian E. and Michael D. Ward. 2010. *Statistical Methodology: Special Issue on Statistical Methods for the Social Sciences*. Elsevier. Volume 8.

5.2. **NSF Grants Received by Jacob M. Montgomery.**

(1) SES-1023762 ($12,000) Doctoral Dissertation Research in Political Science: The Causes, Consequences, and Measurement of Perceived Political Control, 09/2010-08/2011.

Summary of Findings: I developed, evaluated, and validated a measure of perceived political control. I show that the scale is distinct from extant measures, and has superior explanatory power for predicting important behaviors in the future. In work currently underway, I show that actively supporting a winning candidate plays a causal role in shaping these control beliefs.

Broader Impacts: I have presented the research at a conference. In addition, I developed an online platform for applying computer-based psychological testing techniques for more efficient administration of large scale in online surveys to be released for public use.

Publications:

- Jacob M. Montgomery, "An Evolutionary Theory of Democracy: Dynamic Evolutionary Models of American Party Competition with an Empirical Application to the Case of Abortion Policy from 1972-2010", Doctoral Dissertation, Duke University, Sept (2010).

REFERENCES

Abramowitz, Alan I. 2008. "Forecasting the 2008 Presidential Election with the Time-for-Change Model." *PS: Political Science & Politics* 41(4):691–695.

Andriole, Stephen J. and Robert A. Young. 1977. "Toward the Development of an Integrated Crisis Warning System." *International Studies Quarterly* 21(1):107–150.

Armstrong, J. Scott. 2001. Combining Forecasts. In *Principles of Forecasting: A Handbook for Researchers and Practitioners*, ed. J. Scott Armstrong. Norwell, MA: Kluwer Academic Publishers.

Bartels, Larry. 1997. "Specification Uncertainty and Model Averaging." *American Journal of Political Science* 41(2):641–674.

Bartels, Larry M. and John Zaller. 2001. "Presidential Vote Models: A Recount." *PS: Political Science and Politics* 34(1):9–20.

Bennett, D. Scott and Allan C. Stam. 2009. "Revisiting Predictions of War Duration." *Conflict Management and Peace Science* 26(3):256–267.

Brandt, Patrick T., John R. Freeman and Philip A. Schrodt. 2011. "Real Time, Time Series Forecasting of Inter- and Intra-State Political Conflict." *Conflict Management and Peace Science* 28(1):41–64.

Brandt, Patrick T., Michael Colaresi and John R. Freeman. 2008. "The Dynamics of Reciprocity, Accountability, and Credibility." *The Journal of Conflict Resolution* 52(3):343–374.

Brier, Glenn W. 1950. "Verification of Forecasts Expressed in Terms of Probability." *Monthly Weather Review* 78(1):1–3.

Brown, Lloyd B. and Henry W. Chappell. 1999. "Forecasting Presidential Elections using History and Polls." *International Journal of Forecasting* 15(2):127–135.

Campbell, James E. 1992. "Forecasting the Presidential Vote in the States." *American Journal of Political Science* 36(2):386–407.

Campbell, James E. 2008. "The Trial-heat Forecast of the 2008 Presidential Vote: Performance and Value Considerations in an Open-Seat Election." *PS: Political Science & Politics* 41(4):697–701.

Campbell, Jospeh E. and Ken A. Wink. 1990. "Trial-heat Forecasts of the Presidential Vote." *American Politics Research* 18(3):251.

Choucri, Nazli and Thomas W. Robinson, eds. 1978. *Forecasting in International Relations: Theory, Methods, Problems, Prospects*. San Francisco: W.H. Freeman.

Clyde, Merlise. 2003. Model Averaging. In *Subjective and Objective Bayesian Statistics*, ed. S. James Press. 2nd ed. Hoboken, NJ: Wiley-Interscience chapter 13, pp. 320–335.

Clyde, Merlise and Edward I. George. 2004. "Model Uncertainty." *Statistical Science* 19(1):81–94.

Cuzàn, Alfred G. and Charles M. Bundrick. 2008. "Forecasting the 2008 Presidential Election: A Challenge for the Fiscal Model." *PS: Political Science & Politics* 41(4):717–722.

Davies, John L. and Ted Robert Gurr. 1998. *Preventive Measures: Building Risk Assessment and Crisis Early Warning Systems*. Lanham, Md: Rowman & Littlefield Publishers.

De Marchi, Scott, Christopher Gelpi and Jeffrey D. Grynaviski. 2004. "Untangling Neural Nets." *American Political Science Review* 98(2):371–378.

de Mesquita, Bruce Bueno. 2002. *Predicting Politics*. Columbus, OH: Ohio State University Press.

Draper, David. 1995. "Assessment and Propagation of Model Uncertainty." *Journal of the Royal Statistical Society. Series B (Methodological)* 57(1):45–97.

Enders, Walter and Todd Sandler. 2005. "After 9/11: Is it All Different Now?" *The Journal of Conflict Resolution* 49(2):259–277.

Erikson, Robert S. and Christopher Wlezien. 2008. "Leading Economic Indicators, the Polls, and the Presidential Vote." *PS: Political Science & Politics* 41(4):703–707.

Fair, Ray C. 1978. "The Effect of Economic Events on Votes For President." *The Review of Economics and Statistics* 60(2):159–173.

Fair, Ray C. 2010. "Presidential and Congressional Vote-Share Equations: November 2010 Update." Working Paper, Yale University. http://fairmodel.econ.yale.edu/RAYFAIR/PDF/2010C.pdf (accessed June 07, 2011).

Fair, Ray C. 2011. "Vote-Share Equations: November 2010 Update." Working Paper, Yale University. http://fairmodel.econ.yale.edu/vote2012/index2.htm (accessed March 07, 2011).

Fearon, James D. and David D. Laitin. 2003. "Ethnicity, Insurgency and Civil War." *American Political Science Review* 97(1):75–90.

Feder, Stanley A. 2002. "Forecasting For Policy Making in the Post-Cold War Period." *Annual Review of Political Science* 5:111–125.

Freeman, John R. and Brian L. Job. 1979. "Scientific Forecasts in International Relations: Problems of Definition and Epistemology." *International Studies Quarterly* 23(1):113–143.

Geer, John and Richard R. Lau. 2006. "Filling in the Blanks: A New Method for Estimating Campaign Effects." *British Journal of Political Science* 36(2):269–290.

Gill, Jeff. 2004. "Introduction to the Special Issue." *Politi* 12(4):647–674.

Gleditsch, Kristian Skrede and Michael D. Ward. 2010. "Contentious Issues and Forecasting Interstate Disputes." Presented to the 2010 Annual Meeting of the International Studies Association, New Orleans, LA.

Graefe, Andreas, Aldfred G. Cuzan, Randal J. Jones and J. Scott Armstrong. 2010. "Combining Forecasts for U.S. Presidential Elections: The PollyVote." Working Paper. http://dl.dropbox.com/u/3662406/Articles/Graefe_et_al_Combining.pdf (accessed May 15, 2011).

Greenhill, Brian, Michael D. Ward and Audrey Sacks. 2011. "The Separation Plot: A New Visual Method for Evaluating the Fit of Binary Data." *American Journal of Political Science* in press.

Gurr, Ted Robert and Mark Irving Lichbach. 1986. "Forecasting Internal Conflict: A Competitive Evaluation of Empirical Theories." *Comparative Political Studies* 19(3):3–38.

Hamill, Thomas S., Jeffrey S. Whitaker and X. Wei. 2004. "Ensemble Reforecasting: Improving Medium-range Forecast Skill Using Retrospective Forecasts." *Monthly Weather Review* 132(6):1434 – 1447.

Hersbach, Hans. 2000. "Decomposition of the Continuous Ranked Probability Score for Ensemble Prediction Systems." *Weather and Forecasting* 15(5):559–570.

Hibbs, Douglas. 2011. "The Bread and Peace Model Applied to the 2008 US Presidential Election." http://douglas-hibbs.com/Election2008/2008Election-MainPage.htm (accessed March 08, 2011).

Hibbs, Douglas A. 2000. "Bread and Peace Voting in US presidential Elections." *Public Choice* 104(1):149–180.

Hildebrand, David K., James D. Laing and Howard Rosenthal. 1976. "Prediction Analysis in Political Research." *The American Political Science Review* 70(2):509–535.

Hoeting, Jennifer A., David Madigan, Adrian E. Raftery and Christopher T. Volinsky. 1999. "Bayesian Model Averaging: A Tutorial." *Statistical science* 14(4):382–417.

Hoff, Peter D. 2007. "Extending the Rank LIkelihood for Semiparametric Copula Estimation." *Annals of Applied Statistics* 1(1):265–283.

Holbrook, Thomas M. 2008. "Incumbency, National Conditions, and the 2008 Presidential Election." *PS: Political Science & Politics* 41(4):709–712.

Imai, Kosuke and Gary King. 2004. "Did Illegal Overseas Absentee Ballots Decide the 2000 US Presidential Election?" *Perspectives on Politics* 2(3):537–549.

Jerome, Bruno, Veronique Jerome and Michael Lewis-Beck. 1999. "Polls Fail in France: Forecasts of the 1997 Legislative Election." *International Journal of Forecasting* 15(2):163–174.

King, Gary and Langche Zeng. 2001. "Improving Forecasts of State Failure." *World Politics* 53(4):623–658.

Krause, George A. 1997. "Voters, Information Heterogeneity, and the Dynamics of Aggregate Economic Expectations." *American Journal of Political Science* 41(4):1170–1200.

Leblang, David and Shanker Satyanath. 2006. "Institutions, Expectations, and Currency Crises." *International Organization* 60(1):245–262.

Lewis-Beck, Michael S. 2005. "Election Forecasting: Principles and Practice." *The British Journal of Politics & International Relations* 7(2):145–164.

Lewis-Beck, Michael S. and Charles Tien. 2008. "The Job of President and the Jobs Model Forecast: Obama for '08?" *PS: Political Science & Politics* 41(4):687–690.

Lockerbie, Brad. 2008. "Election Forecasting: The Future of the Presidency and the House." *PS: Political Science & Politics* 41(4):713–716.

Madigan, David and Adrian E. Raftery. 1994. "Model Selection and Accounting for Model Uncertainty in Graphical Models Using Occam's Window." *Journal of the American Statistical Association* 89(428):1535–1546.

Marshall, Monty G., Keith Jaggers and Ted Robert Gurr. 2009. "Polity IV Project: Political Regime Characteristics and Transition 1800-2007." CIDCM: University of Maryland, MD.

Martin, Andrew D., Kevin M. Quinn, Theodore W. Ruger and Pauline T. Kim. 2004. "Competing Approaches to Predicting Supreme Court Decision Making." *Perspectives on Politics* 2(2):761–767.

Mesquita, Bruce Bueno De. 2011. "A New Model for Predicting Policy Choices: Preliminary Tests." *Conflict Management and Peace Science* 28(1):65–85.

Montgomery, Jacob M. and Brendan Nyhan. 2010. "Bayesian Model Averaging: Theoretical Developments and Practical Applications." *Political Analysis* 18(2):245–270.

Norpoth, Helmut. 2008. "On the Razor's Edge: The Forecast of the Primary Model." *PS: Political Science & Politics* 41(4):683–686.

O'Brien, Sean P. 2002. "Anticipating the Good, the Bad, and the Ugly: An Early Warning Approach to Conflict and Instability Analysis." *Journal of Conflict Resolution* 46(6):791–811.

O'Brien, Sean P. 2010. "Crisis Early Warning and Decision Support: Contemporary Approaches and Thoughts on Future Research." *International Studies Review* 12(1):87–104.

Pevehouse, Jon C. and Joshua S. Goldstein. 1999. "Serbian Compliance or Defiance in Kosovo? Statistical Analysis and Real-Time Predictions." *The Journal of Conflict Resolution* 43(4):538–546.

Pitt, Mark A., In Jae Myung and Shaobo Zhang. 2002. "Toward a Method of Selecting Among Computational Models of Cognition." *Psychological Review* 109(3):472–491.

Pitt, Mark A. and Jae Myung. 2002. "When a Good Fit Can Be Bad." *TRENDS in Cognitive Sciences* 6(10):421–425.

Raftery, Adrian E. 1995. "Bayesian Model Selection in Social Research." *Sociological Methodology* 25(1):111–163.

Raftery, Adrian E., Tilmann Gneiting, Fadoua Balabdaoui and Michael Polakowski. 2005. "Using Bayesian Model Averaging to Calibrate Forecast Ensembles." *Monthly Weather Review* 133(5):1155–1174.

Raftery, Adrian E. and Yingye Zheng. 2003. "Long-run Performance of Bayesian Model Averaging." *Journal of the American Statistical Association* 98(464):931–938.

Ruger, Theadore .W., Pauline T. Kim, Andrew D. Martin and Kevin M. Quinn. 2004. "The Supreme Court Forecasting Project: Legal and Political Science Approaches to Predicting Supreme Court Decisionmaking." *Columbia Law Review* 104(4):1150–1210.

Schneider, Gerald, Nils Petter Gleditsch and Sabine Carey. 2011. "Forecasting in International Relations: One Quest, Three Approaches." *Conflict Management and Peace Science* 28(1):5–14.

Schrodt, Philip A. and Deborah J. Gerner. 2000. "Using Cluster Analysis to Derive Early Warning Indicators for Political Change in the Middle East, 1979-1996." *American Political Science Review* 94(4):803–818.

Singer, J. David and Michael D. Wallace. 1979. *To Augur Well: Early Warning Indicators in World Politics*. Beverly Hills, CA: Sage Publications.

Sloughter, J. McLean, Adrian E. Raftery, Tilmann Gneiting and Chris Fraley. 2007. "Probabilistic Quantitative Precipitation Forecasting Using Bayesian Model Averaging." *Monthly Weather Review* 135(9):3209–3220.

Sloughter, J. McLean, Tilmann Gneiting and Adrian E. Raftery. 2010. "Probabilistic Wind Speed Forecasting Using Ensembles and Bayesian Model Averaging." *Journal of the American Statistical Association* 105(489):25–35.

Vincent, Jack E. 1980. "Scientific Prediction versus Crystal Ball Gazing: Can the Unknown be Known?" *International Studies Quarterly* 24(3):450–454.

Vrugt, Jasper A., Cees G.H. Diks and Martyn P. Clark. 2008. "Ensemble Bayesian Model Averaging Using Markov Chain Monte Carlo Sampling." *Environmental Fluid Mechanics* 8(5):579–595.

Ward, Michael D., Randolph M. Siverson and Xun Cao. 2007. "Disputes, Democracies, and Dependencies: A Reexamination of the Kantian Peace." *American Journal of Political Science* 51(3):583–601.

Whiteley, Paul F. 2005. "Forecasting Seats from Votes in British General Elections." *The British Journal of Politics & International Relations* 7(2):165–173.

MICHAEL D. WARD

Professional Preparation.

B.A. (Honors), Indiana University, 1970.

Ph. D., Political Science, Northwestern University, 1977.

Gordon Scott Fulcher Post Doctoral Research Associate & Visiting Assistant Professor, Department of Political Science, Northwestern University, 1977-1980.

Appointments.

Professor of Political Science, Duke University, 2009–present.

Professor of Political Science, University of Washington, 1997–2009.

Chaire Municipale & Professeur, Faculté des Sciences Économiques, Université Pierre Mendès France, Grenoble, France, 2000-2005.

Professeur Invitée et Emerité, Faculté des Sciences Économiques, Université Pierre Mendès France, Grenoble, France, 1990-1991.

Professor of Political Science, University of Colorado, 1982-1998.

Research Scientist, Wissenschaftszentrum Berlin, International Institute for Comparative Social Research, 1980-1982.

Gordon Scott Fulcher Post Doctoral Research Associate & Visiting Assistant Professor, Department of Political Science, Northwestern University, 1977-1980.

Publications.

Related to Proposed Activities.

1 Weidmann, Nils B. and Michael D. Ward. 2010. "Predicting Conflict in Space and Time". *Journal of Conflict Resolution* 54(6): 883-910.

2 Ward, Michael D., Brian D. Greenhill and Kristin M. Bakke. 2010. "The Perils of Policy by P-Value: Predicting Civil Conflicts". *Journal of Peace Research* 48(4):3673-375. (Winner of best article in journal during 2010).

3 Ward, Michael D. and Kristian Skrede Gleditsch. 2008. *Spatial Regression Models*. Thousand Oaks, California and London: Sage Publishers

4 Greenhill, Brian D., Michael D. Ward and Audrey Sacks. 2011. "The Separation Plot: A New Visual Method for Evaluating the Predictive Power of Binary Models". *American Journal of Political Science* in press.

5 Ward, Michael D. and Kristian Skrede Gleditsch. 2002. "Location, Location, Location: An MCMC Approach to Modeling the Spatial Context of War and Peace". *Political Analysis* 10(3):244–260.

Five Other Significant Publications.

1 Ward, Michael D., Randolph M. Siverson and Xun Cao. 2007. "Disputes, Democracies, and Dependencies: A Re-examination of the Kantian Peace". *American Journal of Political Science* 51(3):583-601

2 Ward, Michael D. and Kristian Skrede Gleditsch. 1998. "Democratizing for Peace". *American Political Science Review* 92(1):51–61

3 Ward, Michael D. and David R Davis. 1992. "Sizing up the Peace Dividend: Economic Growth and Military Spending in the United States, 1948–1996". *American Political Science Review* 86(3):748–758.

4 Mintz, Alex and Michael D. Ward. 1989. "The Political Economy of Military Spending in Israel". *American Political Science Review* 83(2):521–533.

5 Ward, Michael D. 1984. "Differential Paths to Parity: A Study of the Contemporary Arms Race". *American Political Science Review* 78(2):297–317.

Synergistic Activities. While on the executive council of the Center for Statistics and the Social Sciences, I helped to develop a university wide curriculum in applied statistical methods that would be available to social science departments across the university. We also created new fields of applied statistics within a total of five departments; our first was in the department of Political Science. These served to broaden and deepen training in statistical methods for social science graduate students at the University of Washington.

I have also organized several conferences aimed at helping scholars in the social sciences to apply new approaches to the study of social phenomenon. One of these was a conference at the University of Washington on Spatial and Network approaches to studying political phenomenon, which was sponsored in part by the Center for Statistics and the Social Sciences, but also funded in part by a grant from the Center for Spatially Informed Social Sciences.

In 2010, I also organized and was the local host for the 2010 Duke Conference on Networks in Political Science, an international and multidisciplinary gathering of about 200 scholars and practitioners of network science.

Collaborators.

Collaborators and co-Editors in the past 48 months. John Ahlquist, Kristin Bakke, Xun Cao, Kristin Gleditsch, Brian Greenhill, Peter Hoff, Aseem Prakash, Ian Lustic, Nils Metternich, Clionadh Raleigh, Adrian Raftery, Phil Schrodt, Stephen Shellman, Randolph Siverson, Katherine Stovall, & Nils Weidmann.

Graduate Advisors and Postdoctoral Sponsors.

My advisors: James A. Caporaso (University of Washington), Harold Guetzkow (RIP), Karl W. Deutsch (RIP)

My advisees: John Ahlquist (University of Wisconsin), Xun Cao (Penn State University), Brian Greenhill (Dartmouth University), Nils Weidmann (Yale University), Umut Aydin (Boğaziçi University), Christian Breunig (University of Toronto), Kristian Gleditsch (Essex University), David Brown (University of Colorado), Michael Shin (UCLA)

JACOB M. MONTGOMERY

Professional Preparation.

> B.A. (*Summa Cum Laude*), Wake Forest University, 2002.
> M.S., Statistical Science, Duke University, 2009.
> Ph. D., Political Science, Duke University, 2011.

Appointments.

> Assistant Professor of Political Science, Washington University in St. Louis, 2011–present.

Publications.

Related to Proposed Activities.

> 1 Montgomery, Jacob M. and Brendan Nyhan. 2010. "Bayesian Model Averaging: Theoretical Developments and Practical Applications". *Political Analysis* 18(2): 245-270.
> 2 Montgomery, Jacob M., Alexandra Cooper, Jerome Reiter and Shuo Guan. 2008. "Non-Response Bias on Dimensions of Political Activity Amongst Political Elites". *International Journal of Public Opinion Research* 20(4): 223-231.

Other Significant Publications.

> 1 Aldrich, John H., Jacob M. Montgomery and Wendy Wood. Forthcoming. "Turnout as a Habit".*Political Behavior*. `http://www.springerlink.com/content/4550l465ux177435,`
> 2 Montgomery, Jacob M., Kristie Long Foley and Mark Wolfson. 2006. "Enforcing the Minimum Drinking Age: State, Local, and Agency Characteristics Associated With Compliance Checks and Cops in Shops Programs". *Addiction* 101(2): 223-231.

Synergistic Activities. From 2008-2010, I helped found and organize a weekly methods workshop involving faculty and graduate students at both Duke University and the University of North Carolina.

Collaborators.

Collaborators and co-Editors in the past 48 months. John H. Aldrich (Duke), Alexandra Cooper (Duke), Christophe DeSante (Duke), Melanie Freeze (Duke), Shuo Guan, David Neal (University of Southern California – Psychology), Brendan Nyhan (Dartmouth), Jerome Reiter (Duke – Statistical Science), David Sparks (Duke), Wendy Wood (University of Southern California – Psychology).

Graduate Advisors and Postdoctoral Sponsors.

> Ph.D. Dissertation Advisors: John H. Aldrich (Duke), Nancy Burns (University of Michigan), Michael C. Munger (Duke), David Rohde (Duke)
> M.S. Thesis Advisors: Jerome Reiter (Duke – Statistical Science), David Dunson (Duke – Statistical Science)

DATA MANAGEMENT PLAN

Data and programs from this research will be deposited in a Dataverse archive at `thedata.org` within 12 months of the end of the project.

The PI, Michael Ward, is budgeted for 0.5 months in each year. The faculty benefit rate for Ward is XX%.

The PIs ask for %XX,XXX to cover the cost of a graduate assistant at Duke University budgeted to work at 50% effort for one year (June, 2012 - May, 2013). The graduate student benefit rate is XX.X%. The combined cost for graduate student assistance will be $XX,XXX.

The PIs ask for $2,500 for travel. This money is expected to cover 1 trip for in-person collaborative research in St. Louis, MO. In addition, the PI and one graduate student will attend the yearly Political Methodology Conference.

Finally, Duke University's current indirect cost rate with DHHS is 6X.X%.

BUDGET JUSTIFICATION (WASHINGTON UNIVERSITY IN ST. LOUIS)

The CO-PI, Jacob M. Montgomery, is budgeted for 1 months in each year. The faculty benefit rate for Montgomery is XX%. The combined cost will be $25,000.

The PIs ask for %XX,XXX to cover the cost of a graduate assistant at Washington University in St. Louis budgeted to work at 50% effort for one year (June, 2012 - May, 2013). The graduate student benefit rate is XX.X%. The combined cost for graduate student assistance will be $XX,XXX.

The PIs ask for $2,500 for travel. This money is expected to cover 2 trips for in-person collaborative research in Durham, NC. In addition, Montgomery will attend the yearly Political Methodology Conference.

Finally, Washington University's current indirect cost rate with DHHS is 5X.X%.

TODO

(1) Florian to review references to ensure that they include: article title, journal title, volume number, page numbers, authors, and year of publication.
(2) Budget (one for each school)
(3) Budget justification (One for each school)
(4) Facilities, Equipment and Other Resources (one for each school)
(5) Make all of this conform to the "collaborative proposals" guidelines reproduced belo.
(6) Get the Fastlane "cover sheet" filled out correctly (and identically in both locations).

4. Collaborative Proposals (pg. II-23

A collaborative proposal is one in which investigators from two or more organizations wish to collaborate on a unified research project. Collaborative proposals may be submitted to NSF in one of two methods: as a single proposal, in which a single award is being requested (with subawards administered by the lead organization); or by simultaneous submission of proposals from different organizations, with each organization requesting a separate award. In either case, the lead organizations proposal must contain all of the requisite sections as a single package to be provided to reviewers (that will happen automatically when procedures below are followed). All collaborative proposals must clearly describe the roles to be played by the other organizations, specify the managerial arrangements, and explain the advantages of the multi-organizational effort within the Project Description. PIs are strongly encouraged to contact the cognizant NSF Program Officer prior to submission of a collaborative proposal.

a. Submission of a collaborative proposal from one organization The single proposal method allows investigators from two or more organizations who have developed an integrated research project to submit a single, focused proposal. A single investigator bears primary responsibility for the administration of the grant and discussions with NSF, and, at the discretion of the organizations involved, investigators from any of the participating organizations may be designated as co-PIs. Please note, however, that if awarded, a single award would be made to the submitting organization, with any collaborators listed as subawards.

If a proposed subaward includes funding to support postdoctoral researchers, the mentoring activities to be provided for such individuals must be incorporated in the supplemental mentoring plan outlined in GPG Chapter II.C.2j.

By submission of the proposal, the organization has determined that the proposed activity is administratively manageable. NSF may request a revised proposal, however, if it considers that the project is so complex that it will be too difficult to review or administer as presented. (See GPG Chapter II.C.2g.(vi)(e) for additional instructions on preparation of this type of proposal.)

b. Submission of a collaborative proposal from multiple organizations

In many instances, simultaneous submission of proposals that contain the same Project Description from each organization might be appropriate. For these proposals, the project title must begin with the words "Collaborative Research:". The lead organization's submission will include a Cover Sheet, Project Summary, Project Description, References Cited, Biographical Sketches, Budgets and Budget Justification, Current and Pending support, and Facilities, Equipment and Other Resources for their organization. If applicable, the lead organizations submission also must include a supplemental mentoring plan that must not exceed one page, and that addresses the mentoring activities to be provided for all postdoctoral researchers supported under the entire collaborative project. See GPG Chapter II.C.2.j for additional guidance on mentoring and data management plan requirements for collaborative proposals. Non-lead organization submissions will include all of the above for their organization except the project summary, project description, and references

cited which are the same for all collaborating organizations. FastLane will combine the proposal submission for printing or electronic viewing.

To submit the collaborative proposal, the following process must be completed:

(i) Each non-lead organization must assign their proposal a proposal PIN. This proposal PIN and the temporary proposal ID generated by FastLane when the non-lead proposal is created must be provided to the lead organization before the lead organization submits its proposal to NSF.

(ii) The lead organization must then enter each non-lead organization(s) proposal PIN and temporary proposal ID into the FastLane lead proposal by using the "Link Collaborative Proposals" option found on the FastLane "Form Preparation" screen. Given that such separately submitted proposals constitute a single proposal submission to NSF, it is imperative that the proposals be submitted within a reasonable timeframe to one another.

(iii) All components of the collaborative proposal must meet any established deadline, and, failure to do so may result in the entire collaborative proposal being returned without review.