Statewide Trial-Heat Polls and the 2000 Presidential Election: A Forecast Model*

Jay A. DeSart, *Florida Atlantic University*

Thomas M. Holbrook, *University of Wisconsin–Milwaukee*

Objective. We provide an examination and update of a presidential election forecasting model that we have previously developed to predict state-level presidential election outcomes. *Method.* Our model consists of September statewide trial-heat polls in 1992, 1996, and 2000 along with a prior vote variable. We use this model to generate predictions for both state-level and national-level outcomes. *Results.* Although our model generated reasonably accurate point estimates of state-level outcomes in 2000, it still incorrectly predicted a Gore victory in 2000. *Conclusions.* We discuss possible explanations for the 2000 misprediction and present updated coefficients to be used to generate a forecast for the 2004 election.

"Closest election in 40 years" was the mantra recited time and again by pundits observing the 2000 presidential election campaign. Little did they know that the extremely close outcomes in a number of states, coupled with the Florida recount controversy, would prove that prophecy to be quite true indeed. With Al Gore garnering a popular vote margin of roughly one-half of 1 percent of the over 105 million votes cast, and George W. Bush winning the electoral vote with one more than the minimum number needed to secure victory, the 2000 presidential election will be forever remembered as one of the most closely contested in American history.

Of course, the incredible closeness of the outcome makes the 2000 election unique not just because it was close but because conventional wisdom held that it should not have been close at all. Given the seemingly robust economy and President Clinton's relatively high approval rating, Al Gore *should* have been able to win with relative ease. Indeed, a number of academic forecasting models, which were generally based on this conventional wisdom, predicted that Gore would win, and some predicted he would do so by a fairly wide margin (Abramowitz, 2001; Campbell, 2001a; Holbrook, 2001; Lewis-Beck and Tien, 2001; Lockerbie, 2001; Norpoth, 2001; Wlezien and Erikson, 2001).

*Direct correspondence to Jay A. DeSart, Department of Political Science, Florida Atlantic University, 5353 Parkside Drive, Jupiter, FL 33436 ⟨jdesart@fau.edu⟩. The authors will share all data and coding information with those wishing to replicate this study.

SOCIAL SCIENCE QUARTERLY, Volume 84, Number 3, September 2003

In addition, the results of the 2000 election also highlight the important role that states play in the presidential selection process. Votes are cast and counted at the state level, and state laws dictate the rules by which that process is to take place. As the world's eyes were turned to the drama unfolding in Florida, it was made resoundingly clear that the states are the key players in the process by which Americans select their president.

In this article, we take a close look at the predictability of the 2000 presidential election not on the basis of a "national" model but on the basis of a model for forecasting presidential outcomes in the states. We do so by expanding on a forecasting model initially applied to the 1992 and 1996 elections (Holbrook and DeSart, 1999). Ultimately, we find that although our model incorrectly predicted a Gore victory, statewide trial-heat polls taken in September did reasonably well in predicting the outcome of the 2000 election. Furthermore, by a wide range of criteria the polls generally did a better job of predicting the 2000 election than they did in the two previous elections.

A State-Level Forecasting Model

The forecasting model we use here is quite parsimonious and was relatively accurate in predicting the state-level outcomes of both the 1992 and 1996 presidential elections. Equation 1 represents that model.

$$\text{VOTE}_i = \alpha + \beta_1(\text{POLL})_i + \beta_2(\text{PRIOR VOTE})_i \qquad (1)$$

This simple model consists of two simple predictor variables: the Democratic candidate's average share of the two-party vote in statewide trial-heat polls taken in September (POLL); and the average Democratic share of the two-party popular vote in each state across the two previous presidential elections (PRIOR VOTE). The use of the trial-heat polls not only picks up patterns of long-term party support, but also accounts for election-specific factors that may influence outcomes both nationally and at the state level. And while the September poll variable is expected to be a strong predictor of the eventual outcome, the prior vote should bring additional predictive power to the model, as it accounts for long-term voting trends in each state above and beyond those expressed in contemporaneous polls.

The Polls

The data for our primary variable of interest, statewide presidential trial-heat polls, are taken from the "Poll Track" at ⟨NationalJournal.com⟩.[1] National Journal maintains a fairly comprehensive archive of public opinion

[1]This website is the direct descendant of the now defunct ⟨PoliticsNow.com⟩ website, which was the primary data source for the 1992 and 1996 analysis of this model.

data on its website, including a collection of both national and statewide presidential trial-heat polls. Of particular interest to us are the statewide polls conducted during the month of September.

As shown in Figure 1, there is a very strong correlation between the standings of the major party candidates in the polls in September and how they ultimately fared in the election. A Pearson's *r* of 0.96 between the trial-heat polls and election outcomes suggests that the polls had the potential to be fairly powerful predictive tools in 2000. In the data for 1992 and 1996 we found correlations of 0.81 and 0.91, respectively (Holbrook and DeSart, 1999). This illustrates that September polls were slightly more indicative of the eventual outcome in 2000 than they were in either 1992 or 1996.

However, this is only a piece of the picture. Although the polls' point estimates are closely correlated with the eventual results, a true test of their predictive ability comes in terms of the extent to which they can correctly project the ultimate winner of the election in each state. The analysis presented in Table 1 demonstrates that both candidates eventually went on to win every state in which the September polls showed them to be the clear leader. Of course, one of the unique features of the 2000 election was the fact that a number of states had extremely tight races. This is demonstrated by the number of states in Table 1 in which no candidate had a significant lead in the polls. There were 20 states in which the margin between the

FIGURE 1

September Statewide Polls and Actual Results 2000 Presidential Election.

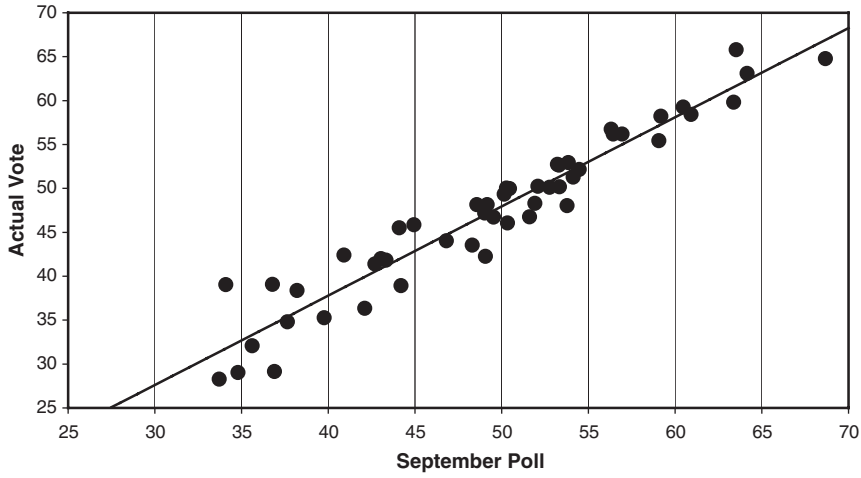

NOTE: September poll variable is Al Gore's average share of the two-party total of the polls conducted in each state in the month of September.

Pearson's $r = 0.96$; $N = 50$; $R^2 = 0.922$; $SE^{y}/_{x} = 2.575$.

TABLE 1

September Polls as Predictors of State-Level Presidential Election Outcomes

	September Poll Status		
Actual Winner	Bush Leading	Toss-Up	Gore Leading
Bush	Alabama Alaska Colorado Idaho Indiana Kansas Mississippi Montana Nebraska North Dakota Oklahoma South Carolina South Dakota Texas Utah Virginia Wyoming *N* = 17	Arizona Arkansas Florida Georgia Kentucky Louisiana Missouri Nevada New Hampshire North Carolina Ohio Tennessee West Virginia *N* = 13	—
Gore	—	Iowa Michigan New Mexico Oregon Pennsylvania Washington Wisconsin *N* = 7	California Connecticut Delaware Hawaii Illinois Maine Maryland Massachusetts Minnesota New Jersey New York Rhode Island Vermont *N* = 13

NOTE: September Poll variable indicates which candidate's average poll tally in each state was higher during the month of September. States were classified as toss-up states if a candidate's lead was within the polls' average margin of error.

Tau-c: 0.814.

candidates' poll standings was within the 95 percent confidence interval (the comparable number for both 1992 and 1996 was 16). Of those 20 states, six were ultimately decided by less than 2 percent of the two-party popular vote, and half of them were decided by less than 4 percent.

If we were to ignore the margin of error and simply project state winners on the basis of the polls' point estimates, the outcomes can be correctly projected in all but six states: Florida, Louisiana, Missouri, New Hampshire,

Tennessee, and West Virginia. Gore held slight leads in each of these states in September, but eventually lost each one by an average of 3.85 percentage points. Of course, in the all-important Electoral College, the electoral votes of any one of these states would have been more than enough to sway the outcome of the election.

The *a Priori* Forecast of the 2000 Election

Ultimately, this preliminary analysis suggests that while the polls conducted in September were closely correlated to the eventual outcome in each state, the projections based on those polls were less than perfect, particularly in light of the extremely close race that eventually came to pass. The question that is of immediate interest to us at this point is how well the model could have predicted the outcome of the 2000 election.

The model in Equation 1 was initially tested for the 1992 and 1996 elections and yielded impressive fit statistics: 90 percent of state outcomes were predicted correctly in 1992 and 92 percent were predicted correctly in 1996; the R^2 statistics were 0.81 and 0.88, respectively. But these estimates were not true forecasting estimates because they were derived based on the relationship between the independent variables and values of the actual outcomes during those years. In other words, they are in-sample post-dictions rather than true forecasts. But pooling the data from 1992 and 1996 allowed us to generate coefficients that could be used to provide a true out-of-sample, *a priori* forecast of the state outcomes for the 2000 election. Those coefficients are found in the first column of Table 2 and are reproduced in the following equation:

$$\text{VOTE}_i = -0.52 + 0.61(\text{POLL})_i + 0.43(\text{PRIOR VOTE})_i \quad (2)$$

Using this formula, we generated forecasts of the 2000 election and present measures of the accuracy of those forecasts in Table 2. The forecasts from the out-of-sample analysis are not as accurate as the results reported for the in-sample analyses of the 1992 and 1996 elections, largely because the actual outcomes from the 2000 election are not used to derive the coefficients.[2] In other words, these are true out-of-sample forecasts.

The *a priori* forecast of the 2000 election missed the eventual winner in 11 (22 percent) states,[3] each of which was predicted by the model to go to Gore, but was ultimately won by Bush. It is interesting to note that while the mean error presented in Table 2 represents the mean *absolute* error in prediction, taking the absolute value turns out to be unnecessary since the *a*

[2]When the 2000 data are used to generate in-sample estimates, the accuracy of the model rivals the in-sample estimated generated for 1992 and 1996, with an R^2 of 0.93 and 92 percent of all state outcomes called correctly.

[3]The states were Arizona, Arkansas, Florida, Kentucky, Louisiana, Missouri, Nevada, New Hampshire, Ohio, Tennessee, and West Virginia.

TABLE 2

Model Performance of Out-of-Sample Forecasts

	2000	1996	1992
September polls	0.61**	0.83**	0.95**
	(0.05)	(0.04)	(0.05)
Prior vote	0.43**	0.16**	0.13*
	(0.06)	(0.04)	(0.06)
Constant	–0.52	–1.72	–5.78*
	(2.47)	(2.34)	(2.24)
N	100	100	100
R^2	0.83	0.85	0.90
$SE_{y/x}$	3.03	3.11	2.81
Out-of-Sample Error			
Mean absolute error	4.81	2.55	2.93
States correctly predicted	78%	90%	86%
National predictions			
Popular vote	54.26%	53.41%	53.74%
Error	+4.07	−1.25	+0.38
Electoral College	376	338	380
Error	+112	–38	+13

NOTE: Figures represent respective model statistics for each year. 2000 coefficients are those reported in our original analysis. $SE_{y/x}$ is the standard error of the estimate. Mean absolute error is the average of the absolute differences between the predicted and actual values of the dependent variable. States correctly predicted is based on an assignment of states to a predicted "winner" based on the model's point estimations for each state. The national predictions are extrapolations from the state-level predictions. Those predictions, and their reported errors, exclude the District of Columbia. Figures in parentheses are standard errors.

*$p < 0.05$; **$p < 0.01$.

priori model overpredicted Gore's level of support in *every state*. The errors ranged from 0.57 in Rhode Island to 10.76 in Wyoming, with a mean of 4.81. This suggests that Gore did worse in every state than would be expected based on his standing in September polls and on Bill Clinton's performance in the two previous elections. Among other things, this could be indicative of general problems with Gore's national campaign, as has been suggested by some forecasters who relied on national indicators in their models (Campbell, 2001a; Holbrook, 2001; Lockerbie, 2001; Wlezien and Erikson, 2001).

To make more appropriate comparisons of model performance across years we also report the results of out-of-sample estimates for 1992 and 1996 in the second and third columns of coefficients in Table 2. These models were estimated by omitting the observations for each year, respectively, and then generating coefficients based on the observations from the other two years. In doing this we are replicating the method used to generate Equation 2 and the *a priori* forecasts of the 2000 election. The primary difference, of course, is that while the estimates for 1992 and 1996 are out-of-sample

estimates, they are not true forecasts since the slopes for 1992 are based on data from 1996 and 2000, and the slopes for 1996 are based on 1992 and 2000.

As one would expect, these predictions are not quite as accurate as the in-sample predictions for the same years (see earlier discussion). Even so, the "forecasts" for 1992 and 1996 are still better than those for 2000, both in terms of mean absolute error (2.55 and 2.93 for 1996 and 1992, respectively) and states correctly predicted (90 percent and 86 percent for 1996 and 1992, respectively). Ultimately, the 2000 election stands out among the three as apparently the most difficult to predict.

National Outcomes

Although the primary focus of the original model was to predict state-level outcomes, the coefficients in Table 2 can also be used as a basis for predicting national outcomes. To predict the Electoral College votes,[4] one simply needs to use the predicted statewide popular vote to project a winner for each state and award the electoral votes to that candidate.

It is also possible to aggregate the individual state estimates to generate a prediction of the national popular vote as well. Because each state contributes a different amount of votes to the overall nationwide vote tally, simply averaging the predicted vote shares across the 50 states would not generate a reliable national estimate, as sparsely populated states and/or states with relatively low turnout would be overrepresented in the calculation of such an estimate. Therefore, we use a weighting scheme that is based on the total number of votes cast in each individual state in prior elections[5] as a proportion of the total number of votes cast nationwide. To minimize the possible effects of idiosyncrasies in turnout in single elections we use the average of those proportions for each state across the previous two elections. The resulting value thus represents the weight each state typically contributes to the overall nationwide turnout. We can then derive a prediction of Gore's share of the national two-party popular vote by simply summing the weighted state-level predictions (weight * predicted vote share).

On the basis of these extrapolations, the model estimated that Gore would get 54.26 percent of the two-party vote, when he actually received 50.19 percent, for an error of 4.07 percentage points. In addition, the Electoral College estimate of 376 votes for Gore overestimated his actual vote by 112. Does this constitute a lot of error? Yes and no. One thing to bear in mind is

[4]Including the District of Columbia would have biased estimates in the tests of the forecast models due to its outlier status. Therefore, the District of Columbia's three electoral votes are excluded from both predicted and actual electoral vote tallies.

[5]Recall that since we are trying to measure each variable prior to the election, we don't know how many votes will be cast in the current election.

that the 2000 election was a difficult one to call, and most forecasting models were fairly wide of the mark. Compared to other forecasting models, the popular vote estimates provided here are slightly better than average.[6]

In terms of the Electoral College vote, it is difficult to make a relative assessment of the error involved because there are not many Electoral College forecasting models with which we can compare our prediction. Lewis-Beck and Rice (1992) have presented a model that predicts the incumbent party's percentage of the Electoral College vote. Their estimates were based on presidential elections from 1948 to 1988, so results from forecasts based on their reported coefficients are not truly comparable.

To address this, we updated their model with data from 1992 and 1996. Our updated Lewis-Beck and Rice (LBR) model would have forecast Al Gore winning 74.27 percent (400) of the electoral votes in 2000, for an error of 134 electors.[7] Even though our error of 112 electoral votes (an error of approximately 21 percent) represents a slight improvement over the LBR model, it still represents considerable error.

September or October?

One question that might arise from an examination of the results so far concerns the timing of the forecast. Our model relies on data that is available just over one month prior to the election. In that regard, its lead-time is somewhat shorter than most election forecast models.[8]

One explanation for our model's misfire on the 2000 election could depend on the movement in the polls during the month of October. Jones (2001) has demonstrated that national-level trial-heat polls conducted in October generally tend to generate more accurate predictions of the eventual outcome than those taken earlier. Therefore, it is reasonable to expect that state-level polls would exhibit a similar pattern.

However, Frankovic and McDermott (2001) and Campbell (2001c) both report that the national-level polls in the fall of 2000 showed a significant amount of stability. Indeed, in our own analysis of the accuracy of state-level polls (Holbrook and DeSart, forthcoming), we found that the polls in 2000 demonstrated relatively little systematic convergence on the outcome after August.

[6]Campbell's (2001b) summary of eight forecasting models found that the average forecast was for a Gore victory with 55.4 percent of the vote. Three of these models provided a more accurate estimate than our model, and four were less accurate.

[7]Our update of the LBR model is $V = -21.74+6.82G+1.41PP$, where V is percent of the electoral votes, G is the real growth rate in GNP over the first two quarters of the election year, and PP is percent approving of the president's job performance in July of the election year (Gallup Poll). Plugging in 1.88 for G and 59 for PP, the equation yields an estimate of 74.27 percent of the electoral vote, or 399.6 votes.

[8]Most published presidential election forecast models use data that is available by the end of August.

When we replicate some of the results we have presented so far using data from October polls we find a modest improvement in the predictions.[9] Looking at Gore's average share in the two-party vote in October polls shows some evidence that the already close race tightened up and shifted slightly toward Bush in the final month of the campaign.

Four states where Gore held a significant lead in September (Delaware, Illinois, Iowa, and Minnesota) ultimately shifted to toss-up status in October. Gore eventually held on to win these states. Four states that were classified as toss-ups in September polls (Georgia, Kentucky, Louisiana, and West Virginia) moved into the Bush column in October and were eventually won by the Republican.

As Table 3 demonstrates, when we generate new out-of-sample predictions using October data we find that the model does indeed produce generally more accurate predictions. Although the *a priori* estimates of the model using September data generated accurate predictions in only 39 states, the model using October data correctly predicted the outcome in 46 states. All four of the incorrectly predicted states (Arkansas, Florida, Missouri, and New Hampshire) had been classified as toss-up states in both the September and October poll data and were eventually won by Bush.

However, it is interesting to note that even though the model using October data produces more accurate predictions of the 2000 state-level outcomes, it *still* incorrectly predicts a Gore victory at the national level in both the popular vote and Electoral College. Furthermore, Table 3 also shows that the degree of improvement that was generated by the October data in the 2000 prediction was not typical in comparison to the predictions for 1992 and 1996. The improvement of the out-of-sample predictions for those two years was relatively modest by comparison.

In the end, how do we evaluate the model using October data? Clearly, there are two factors that one must consider in evaluating a forecast model: (1) the lead-time it provides in generating predictions, and (2) the accuracy of those predictions. By relying on October data, we generate somewhat more accurate predictions while sacrificing a significant amount of lead-time. In effect, our lead-time has dwindled from just over a month prior to the election to between two and seven days. In our opinion, the improvement we find in the model is not enough to justify such a large sacrifice in lead-time.

[9]Ten states did not have polls that were conducted exclusively in the month of October. Those states were AK, CT, HI, KS, MA, MS, NM, OK, UT, and WY. In those cases we used polls from other time frames. In four states (CT, MS, NM, and OK), we selected polls that were at least partially conducted in the month of October, but were started in late September or ended in early November. There were no polls conducted in Utah between September 12 and November 1, so we used the results from a poll conducted from November 1 through November 4. No poll data after September were available for the remaining states (AK, HI, KS, MA, and WY), so for those states the latest September poll available was used.

TABLE 3

Comparison of Out-of-Sample Model Performance September Polls Versus October Polls

	2000	1996	1992
Mean Absolute Error			
September	4.81	2.55	2.93
October	2.61	1.78	1.58
Improvement	2.20	0.77	1.35
States Correctly Predicted			
September	78%	90%	86%
October	92%	92%	96%
Improvement	14%	2%	10%
Popular Vote Prediction			
September	54.26%	53.41%	53.74
Error	+4.07	–1.25	+0.38
October	51.20%	54.18%	53.77
Error	+1.01	–0.48	+0.41
Improvement	3.06	0.77	(0.03)
Electoral College Prediction			
September	376	338	380
Error	+112	–38	+13
October	310	385	379
Error	+46	+9	+12
Improvement	66	29	1

NOTE: Figures represent a comparison of predictions based on September poll model (from Table 2) to those from a model using October polls. Mean Absolute Error is the average of the absolute differences between the predicted and actual values of the dependent variable. States Correctly Predicted is based on an assignment of states to a predicted "winner" based on the model's point estimations for each state. The national predictions are extrapolations from the state-level predictions. Those predictions, and their reported errors, exclude the District of Columbia.

Updating the Model for 2004

Beyond the close nature of the outcome, one of the difficulties in forecasting state outcomes for the 2000 election is that the coefficients of the *a priori* model are based on data from just two years, 1992 and 1996. Although the estimates were based on 100 observations, those observations were drawn from only two different national contexts. To the extent that each election year is unique and the relative weight assigned to each of the forecasting factors changes somewhat from one year to the next (note the instability of the coefficients across years in Table 2), it is somewhat surprising that the 2000 forecasts were as accurate as they were. Presumably, as more data (years) are added to the *a priori* model, the coefficients will begin to approximate the "typical" effect of previous votes and September poll standing and will provide more accurate predictions.

To that end, we have updated our general forecasting model in Table 4 by polling data from the 1992–2000 elections. As the results indicate, the

TABLE 4

General Forecasting Model for State-Level Presidential Election Outcomes

	Unstandardized Regression Coefficient
September poll	0.85**
	(0.04)
Prior vote	0.20**
	(0.04)
Constant	–3.63
	(1.98)
N	150
R^2	0.85
$SE_{y/x}$	3.10
Mean absolute error	2.36

	Model Performance for Individual Years		
	2000	1996	1992
Mean absolute error	2.14	2.38	2.63
States correctly predicted	92%	94%	88%
National predictions			
Popular vote	50.80%	53.98%	52.99%
Error	+0.61	–0.68	–0.37
Electoral College	286	355	372
Error	+22	–21	+5

NOTE: Standard errors in parentheses.

**$p < 0.01$.

model generates in-sample estimates that fit the data fairly well, with 91 percent of all outcomes estimated correctly. The national-level extrapolations of these in-sample predictions are also fairly close to the mark in each year. Although our model still generates a prediction of a Gore Electoral College victory in 2000, the error of 22 electoral votes (4.11 percent) is a bit more acceptable than the error produced by our *a priori* prediction.[10]

This improvement is understandable, since in-sample forecasts *should* be more accurate than out-of-sample forecasts. As a more suitable basis for comparison, we calculated the average absolute error in the in-sample predictions generated by our update of the LBR model. Our average absolute error of 2.99 percent (roughly 16 electoral votes) in the Electoral College prediction across years is substantially better than the 12.09 percent (approximately 65 electoral votes) achieved by the updated LBR model.

[10]In-sample predictions of the 2000 election using October data still estimated a Gore victory in the Electoral College with 298 electoral votes (an error of 34 votes), but estimated that he would *lose in the popular vote* with only 49.93 percent of the two-party popular vote (an error of 0.26 percentage points).

Ultimately, we are fairly pleased with the performance of this model after the addition of the data from the 2000 election, and the new set of coefficients from our updated model can be used to generate out-of-sample, *a priori* forecasts of the 2004 state outcomes once the September polling data are available.

Conclusion

The analysis in this article updated our state-level forecasting model (Holbrook and DeSart, 1999). Several interesting findings emerged. First, as with the earlier analysis, the die appears to be fairly well cast by the end of September. In 2000, as in 1992 and 1996, candidates went on to win almost all the states in which they held a lead in the polls in September, even if that lead was within the margin of error. Further, in 2000, both candidates won *every* state in which they held a lead outside the margin of error in the September polls. Second, when data on past voting patterns are added to the mix, the model produced highly accurate in-sample estimates of the eventual outcomes. Third, it was also possible to use the coefficients from the polled 1992 and 1996 model to produce true forecasts for the 2000 outcomes. Although these estimates contained appreciably more error than the in-sample model did, they were relatively accurate in comparison to other forecasting models (see note 6). Finally, the addition of observations from the 2000 elections provides additional data that have been used to update the coefficients of the *a priori* model. These estimates can now be used to forecast the 2004 election. Only time will tell if the addition of the observations from 2000 increases the accuracy of the general model. Ultimately, we believe that as time goes on and more and more data points are added to the analysis, the resulting coefficients will become more stable and the forecasting performance of our model will improve.

REFERENCES

Abramowitz, Alan. 2001. "The Time for Change Model and the 2000 Election." *American Politics Research* 29:279–82.

Campbell, James E. 2001a. "An Evaluation of the Trial-Heat and Economy Forecast of the Presidential Vote in the 2000 Elections." *American Politics Research* 29:289–96.

———. 2001b. "Taking Stock of the Forecasts of the 2000 Presidential Election." *American Politics Research* 29:275–78.

———. 2001c. "The Curious and Close Presidential Campaign of 2000." In William Crotty, ed., *America's Choice 2000*. Boulder, Colo.: Westview.

DeSart, Jay A., and Thomas Holbrook. Forthcoming. "Campaigns, Polls and the States: Assessing the Accuracy of Statewide Presidential Trial-Heat Polls." *Political Research Quarterly*.

Frankovic, Kathleen A., and Monika L. McDermott. 2001. "Public Opinion in the 2000 Election: The Ambivalent Electorate." In Gerald M. Pomper et al., *The Election of 1992: Reports and Interpretations*. New York: Chatham House.

Holbrook, Thomas M. 2001. "(Mis)Reading the Political Tea Leaves." *American Politics Research* 29:297–301.

Holbrook, Thomas M., and Jay A. DeSart. 1999. "Using State Polls to Forecast Presidential Election Outcomes in the American States." *International Journal of Forecasting* 15:137–42.

Jones Jr., Randall J. 2001. *Who Will Be in the White House? Predicting Presidential Elections*. New York: Longman.

Lockerbie, Brad. 2001. "Forecast 2000: An Afterthought." *American Politics Research* 29:307–12.

Lewis-Beck, Michael S., and Tom W. Rice. 1992. *Forecasting Elections*. Washington, D.C.: CQ Press.

Lewis-Beck, Michael S., and Charles Tien. 2001. "Election 2000: How Wrong Was the Forecast?" *American Politics Research* 29:302–06.

Norpoth, Helmut. 2001. "'Of Time and Candidates' Revisited: Too Generous for Al Gore." *American Politics Research* 29:313–19.

Wlezien, Christopher, and Robert S. Erikson. 2001. "After the Election: Our Forecast in Retrospect." *American Politics Research* 29:320–28.